

EdX_Harvard_Capstone_CYO_PIMA Indian Diabetes_ Prediction

Srividhya G Ammanur

11/18/2020

Executive Summary:

This project is for the data science professional certificate capstone, from HarvardX institution through the edX platform. The objective is to construct a machine learning model that learns from data by looking at a broad range of patterns and make inferences that efficiently predicts the outcome without human intervention. The data set utilized for this project is the PIMA Indian Diabetes Data Set.

Data Source:

Original Owners: National Institute of Diabetes and Digestive and Kidney Diseases. **Donor of database:** Vincent Sigillito (vgs@aplcn.apl.jhu.edu), who was part of Applied Physics Laboratory at Johns Hopkins. The Applied Physics Laboratory (APL) is a not-for-profit engineering research and development center founded in 1942 to assist the military with ballistics detonation. The data set was constructed from a larger database by the NIDDK. For this project the data set have been taken from the UCI Repository Of Machine Learning Databases.

The women behind the Dataset- A background:

All participants (768 Observations) in this dataset are women, at least 21 years old of PIMA Indian heritage. PIMA, are North American Indians who traditionally lived along the Gila and Salt rivers in Arizona, U.S., in what was the core area of the prehistoric Hohokam culture. The PIMA, who speak a Uto-Aztecan language and call themselves the “River People,” are usually considered to be the descendants of the Hohokam. The PIMA Indian women have given a great gift to the humanity by volunteering for this research and donating their biological data for the larger good of the society. I express my sincere gratitude for their generosity. William Knowler, an NIH researcher since 1975, who is also recognized as one of the world’s highly cited researchers in clinical medicine, biology and biochemistry, testified before congress, “This study has contributed much to the world’s current understanding of the causes and consequences of Type 2 diabetes and its complications, for which we are indebted to this community”.

Data Ethics:

It is important for the community of data scientists to be aware that all those who make investments in repositories of data whether it is heart disease, breast cancer research, or social media usage are not always the ones to benefit from their use. As Rebecca Lemov and Dan Bouk rightly said, “data are people too”.

When data rises above the geographically defined locality and circumstances it becomes “big” and is fed into computers and algorithms to meet business goals. It therefore becomes obligatory, that we recognize the living people behind the data. Wherever feasible go beyond and create institutions and frameworks the benefit the contributors of data for data Science.

Project Methodology:

- Gathering data
- Preparing the data
- Choosing a model
- Training the model
- Evaluating the model
- Tuning the model
- Make Predictions

Load the dataset

```
setwd("C:\\Users\\agsri\\OneDrive\\Capstone") diabetes <- read.csv(file = 'Pima_Indian_Diabetes.csv')
```

```
setwd("C:\\Users\\agsri\\OneDrive\\Capstone")  
diabetes <- read.csv(file = 'Pima_Indian_Diabetes.csv')
```

Load Libraries

```
## Warning: package 'ggcorrplot' was built under R version 4.0.2
```

```
## -- Attaching packages ----- tidyverse 1.3.0 --
```

```
## v tibble  3.0.1      v dplyr    0.8.5  
## v tidyr   1.0.3      v stringr 1.4.0  
## v readr   1.3.1      v forcats 0.5.0  
## v purrr   0.3.4
```

```
## -- Conflicts ----- tidyverse_conflicts() --  
## x dplyr::filter() masks stats::filter()  
## x dplyr::lag()     masks stats::lag()
```

```
## Warning: package 'GGally' was built under R version 4.0.2
```

```
## Registered S3 method overwritten by 'GGally':  
##   method from  
##   +.gg      ggplot2
```

```
##  
## Attaching package: 'GGally'
```

```
## The following object is masked from 'package:dplyr':
##
##      nasa

## Warning: package 'caret' was built under R version 4.0.2

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
##      lift
```

view the first and last 10 rows of the dataset

```
head(diabetes)
```

```
##      Pregnancies Glucose BloodPressure SkinThickness Insulin  BMI
## 1           6      148           72           35         0 33.6
## 2           1       85           66           29         0 26.6
## 3           8      183           64            0         0 23.3
## 4           1       89           66           23        94 28.1
## 5           0      137           40           35       168 43.1
## 6           5      116           74            0         0 25.6
##      DiabetesPedigreeFunction Age Outcome
## 1                0.627  50         1
## 2                0.351  31         0
## 3                0.672  32         1
## 4                0.167  21         0
## 5                2.288  33         1
## 6                0.201  30         0
```

```
tail(diabetes)
```

```
##      Pregnancies Glucose BloodPressure SkinThickness Insulin  BMI
## 763           9       89           62            0         0 22.5
## 764          10      101           76           48       180 32.9
## 765           2      122           70           27         0 36.8
## 766           5      121           72           23       112 26.2
## 767           1      126           60            0         0 30.1
## 768           1       93           70           31         0 30.4
##      DiabetesPedigreeFunction Age Outcome
## 763                0.142  33         0
## 764                0.171  63         0
## 765                0.340  27         0
## 766                0.245  30         0
## 767                0.349  47         1
## 768                0.315  23         0
```

Summary Statistics of dataset

glance the data set to see if it is in tidy format

```
diabetes %>% as_tibble()
```

```
## # A tibble: 768 x 9
##   Pregnancies Glucose BloodPressure SkinThickness Insulin   BMI
##   <int>      <int>      <int>      <int>      <int> <dbl>
## 1         6      148         72         35         0  33.6
## 2         1       85         66         29         0  26.6
## 3         8      183         64          0         0  23.3
## 4         1       89         66         23        94  28.1
## 5         0      137         40         35       168  43.1
## 6         5      116         74          0         0  25.6
## 7         3       78         50         32        88   31
## 8        10      115          0          0         0  35.3
## 9         2      197         70         45       543  30.5
## 10        8      125         96          0         0   0
## # ... with 758 more rows, and 3 more variables: DiabetesPedigreeFunction <dbl>,
## #   Age <int>, Outcome <int>
```

Summary of the data set

```
summary(diabetes)
```

```
##   Pregnancies      Glucose    BloodPressure    SkinThickness
##   Min.   : 0.000   Min.   : 0.0   Min.   : 0.00   Min.   : 0.00
##   1st Qu.: 1.000   1st Qu.: 99.0   1st Qu.: 62.00   1st Qu.: 0.00
##   Median : 3.000   Median :117.0   Median : 72.00   Median :23.00
##   Mean   : 3.845   Mean   :120.9   Mean   : 69.11   Mean   :20.54
##   3rd Qu.: 6.000   3rd Qu.:140.2   3rd Qu.: 80.00   3rd Qu.:32.00
##   Max.   :17.000   Max.   :199.0   Max.   :122.00   Max.   :99.00
##   Insulin      BMI      DiabetesPedigreeFunction      Age
##   Min.   : 0.0   Min.   : 0.00   Min.   :0.0780   Min.   :21.00
##   1st Qu.: 0.0   1st Qu.:27.30   1st Qu.:0.2437   1st Qu.:24.00
##   Median : 30.5   Median :32.00   Median :0.3725   Median :29.00
##   Mean   : 79.8   Mean   :31.99   Mean   :0.4719   Mean   :33.24
##   3rd Qu.:127.2   3rd Qu.:36.60   3rd Qu.:0.6262   3rd Qu.:41.00
##   Max.   :846.0   Max.   :67.10   Max.   :2.4200   Max.   :81.00
##   Outcome
##   Min.   :0.000
##   1st Qu.:0.000
##   Median :0.000
##   Mean   :0.349
##   3rd Qu.:1.000
##   Max.   :1.000
```

Structure of the data set

```
str(diabetes)
```

```
## 'data.frame': 768 obs. of 9 variables:
## $ Pregnancies : int 6 1 8 1 0 5 3 10 2 8 ...
## $ Glucose : int 148 85 183 89 137 116 78 115 197 125 ...
## $ BloodPressure : int 72 66 64 66 40 74 50 0 70 96 ...
## $ SkinThickness : int 35 29 0 23 35 0 32 0 45 0 ...
## $ Insulin : int 0 0 0 94 168 0 88 0 543 0 ...
## $ BMI : num 33.6 26.6 23.3 28.1 43.1 25.6 31 35.3 30.5 0 ...
## $ DiabetesPedigreeFunction: num 0.627 0.351 0.672 0.167 2.288 ...
## $ Age : int 50 31 32 21 33 30 26 29 53 54 ...
## $ Outcome : int 1 0 1 0 1 0 1 0 1 1 ...
```

Columns in the dataset

```
colnames(diabetes)
```

```
## [1] "Pregnancies" "Glucose"
## [3] "BloodPressure" "SkinThickness"
## [5] "Insulin" "BMI"
## [7] "DiabetesPedigreeFunction" "Age"
## [9] "Outcome"
```

```
#Rows in the dataset
```

```
nrow(diabetes)
```

```
## [1] 768
```

Average age of the women in the data set

```
mean(diabetes$Age)
```

```
## [1] 33.24089
```

```
mean(diabetes$Pregnancies)
```

```
## [1] 3.845052
```

```
range(diabetes$BMI)
```

```
## [1] 0.0 67.1
```

```
median(diabetes$Glucose)
```

```
## [1] 117
```

Check for any null values in the data

```
colSums(is.na(diabetes))
```

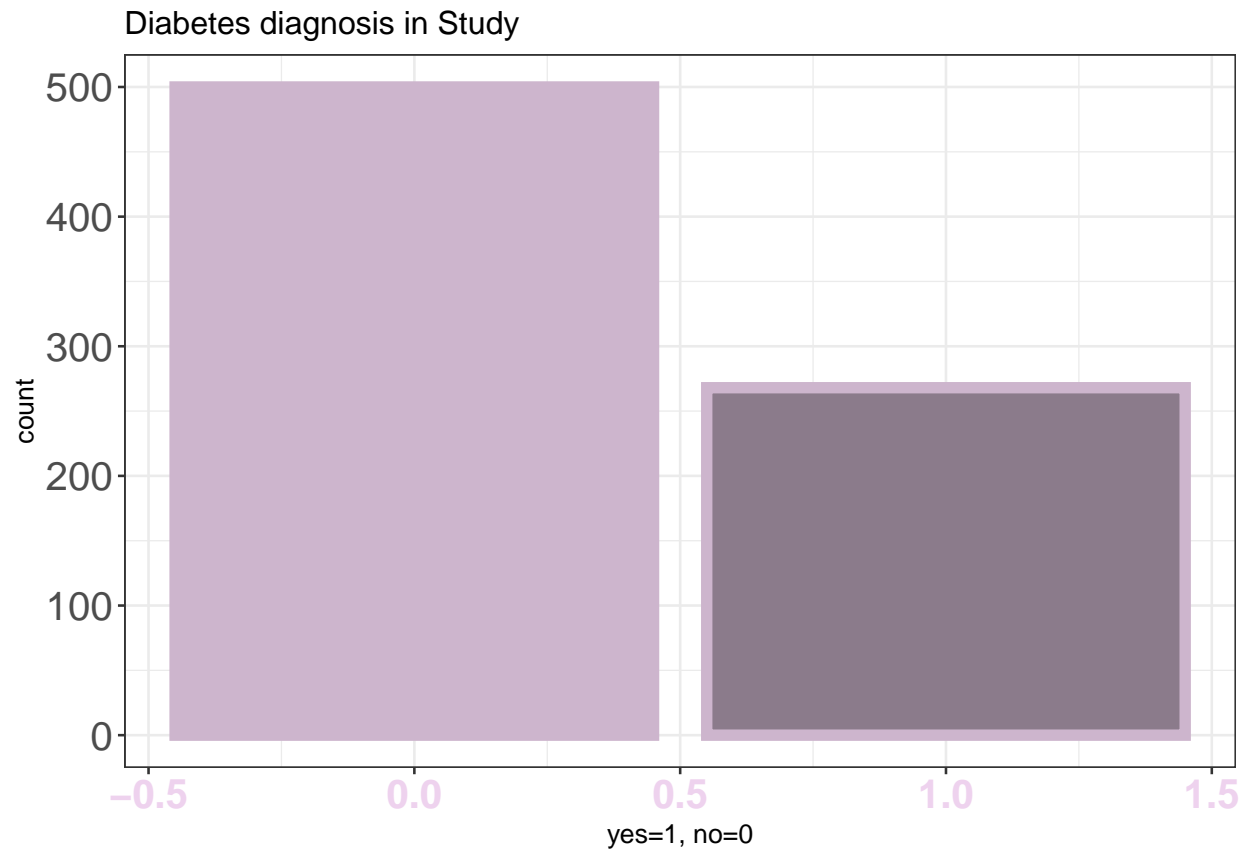
```
##           Pregnancies           Glucose           BloodPressure
##                0                0                0
##       SkinThickness       Insulin                BMI
##                0                0                0
## DiabetesPedigreeFunction      Age           Outcome
##                0                0                0
```

DATA VISUALIZATION

Exploring & understanding the data through visualizations

```
options(repr.plot.width=8, repr.plot.height=7)

ggplot(data = diabetes) +
  geom_bar(stat = "count", mapping = aes(x = Outcome),
          color = 'thistle3', fill = c('thistle3','thistle4'), lwd = 2) +
  labs(x = "yes=1, no=0", title = "Diabetes diagnosis in Study") +
  theme_bw(base_size = 10) +
  theme(axis.text.x = element_text(size = 15,
                                   colour = 'thistle2', face = "bold"),
        axis.text.y = element_text(size = 15))
```



Draw plot on Glucose in blood and diabetic outcome.

```
ggplot(diabetes, aes(x=Glucose, y=Outcome)) +  
  geom_bar(stat="identity", width=.5, fill="tomato3") +  
  labs(title="Ordered Bar Chart",  
        subtitle=" Glucose Vs Diabetic Outcome",  
        caption="source: PIMA Indian dataset") +  
  theme(axis.text.x = element_text(angle=65, vjust=0.6))
```

Ordered Bar Chart

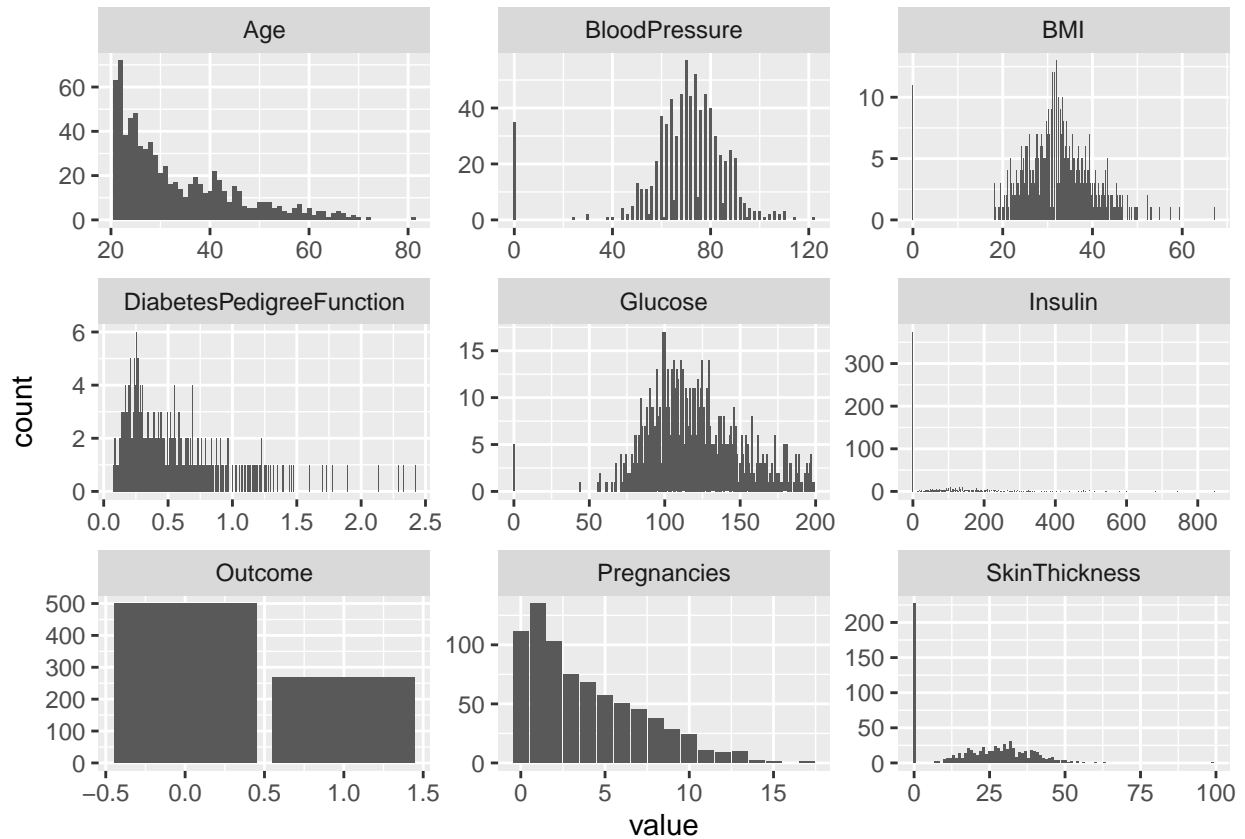
Glucose Vs Diabetic Outcome



source: PIMA Indian dataset

#Histogram of all the variables in the data set.


```
diabetes %>%  
  gather() %>%  
  ggplot(aes(value)) +  
  facet_wrap(~ key, scales = "free") +  
  geom_bar()
```

Plot each variables with diabetic outcome distribution


```
# Plot pregnancy vs outcome distribution
ggplot(data = diabetes,aes(x = Pregnancies)) +
  geom_histogram(binwidth = 0.6,aes(fill = Outcome),position = "dodge") +
  ggtitle("Pregnancies Data Distribution") + ylab("Diabetic OutCome") +
  theme_gray() +
  theme_update(plot.title = element_text(hjust = 0.6))
```

Pregnancies Data Distribution




```
# Plot BMI vs outcome distribution
ggplot(data = diabetes,aes(x = BMI)) +
  geom_histogram(binwidth = 0.6,aes(fill = Outcome),position = "dodge") +
  ggtitle("BMI & Diabetes correlation") + ylab("Diabetic Outcome") +
  theme_gray() +
  theme_update(plot.title = element_text(hjust = 0.6))
```

BMI & Diabetes correlation



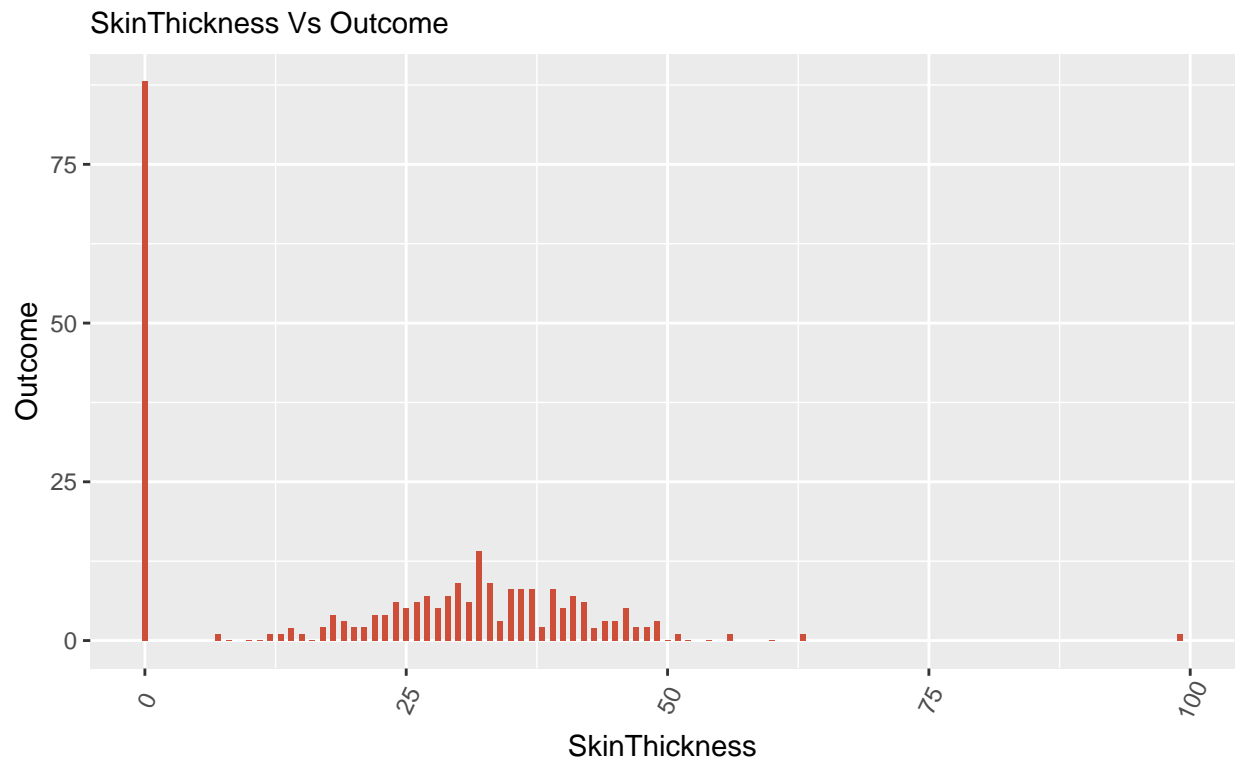
```
# plot Age vs outcome distribution
ggplot(aes(x = Age), data = diabetes) +
  geom_bar(fill='#b47846') + ggtitle("Age & Diabetes outcome distribution")
```

Age & Diabetes outcome distribution



```
# plot Skintickness vs outcome distribution
ggplot(diabetes, aes(x=SkinThickness, y=Outcome)) +
  geom_bar(stat="identity", width=.5, fill="tomato3") +
  labs(title="Ordered Bar Chart",
        subtitle="SkinThickness Vs Outcome",
        caption="source: PIMA Indian Dataset") +
  theme(axis.text.x = element_text(angle=65, vjust=0.6))
```

Ordered Bar Chart




#Proptional table of the diabetes dataset

```
prop.table(table(diabetes$Outcome))
```

```
##
##      0      1
## 0.6510417 0.3489583
```

Correlation between attributes

```
corr<-round(cor(diabetes),1)
ggcorrplot(corr, hc.order = TRUE,
  type = "lower",
  lab = TRUE,
  lab_size = 5,
  method="square",
  colors = c("Yellow", "Grey", "blue"),
  title="Correlogram of Diabetes data",
  ggtheme=theme_bw)
```



#A correlation measures the relationship between two variables. A -ve correlation indicates that if one variable increases the other decreases. A negative correlation exists between pregnancy and insulin and skin thickness. Also age and skin thickness are negatively correlated. Age, pregnancy with diabetes pedigree, are not correlated. Similarly BMI is also not correlated with pregnancy & age. They all show a 0 in the correlogram. The rest of the variables have correlation with values close to 0. The variables show the most correlation are the following:

1. BMI & Diabetes pedigree function
2. Blood Pressure & Insulin
3. Pregnancy & blood pressure
4. Glucose & skin thickness
5. Glucose & pregnancy
6. Glucose & diabetes pedigree function

In order to build a model and train it lets Split data into training set and test data set

```
set.seed(2017)

trainIndex <- createDataPartition(diabetes$Outcome, p = .8, list = FALSE, times = 1)

diabetes$Outcome <- as.factor(diabetes$Outcome)
```

```
diabetes.train <- diabetes[trainIndex,]
diabetes.test <- diabetes[-trainIndex,]
```

Training data Proportion

```
prop.table(table(diabetes.train$Outcome))
```

```
##
##           0           1
## 0.6569106 0.3430894
```

Test data Proportion

```
prop.table(table(diabetes.test$Outcome))
```

```
##
##           0           1
## 0.627451 0.372549
```

Model1: RANDOM FOREST MODEL:

Build a Random Forest Model.

#Individual decisions trees are combined to make a random forest.Each decision tree is the building block of the random forest model. Random forest builds multiple decision trees and merges them together to get a more accurate and stable prediction.

#Fit Random Forest Model in training set data # Train the model using random forest algorithm.

```
control <- trainControl(
  method = "repeatedcv",
  number = 20,
  repeats = 20
)
# Performance Parameters Setting
grid <- expand.grid(mtry = c(3,4,5))

model.Random.Forest <- train(Outcome ~ ., data = diabetes.train,
  method = "rf", tuneGrid = grid,trControl = control)

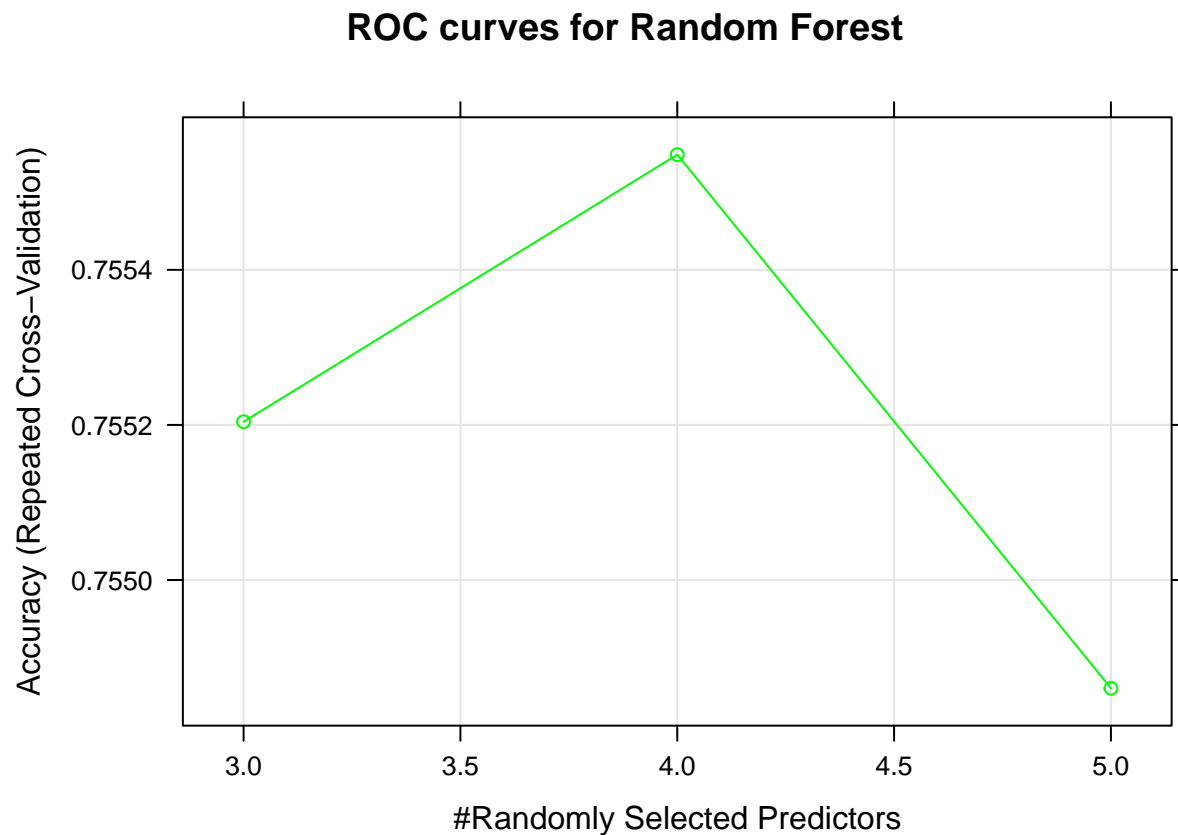
model.Random.Forest
```

```
## Random Forest
##
## 615 samples
```

```
## 8 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (20 fold, repeated 20 times)
## Summary of sample sizes: 584, 585, 583, 585, 584, 585, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
## 3 0.7552041 0.4362251
## 4 0.7555486 0.4380512
## 5 0.7548604 0.4375340
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 4.
```

Plotting ROC curve

```
plot(model.Random.Forest, main = "ROC curves for Random Forest", col='Green')
```



Predict the outcome on the test data

```
predict.Random.Forest <- predict(model.Random.Forest,diabetes.test)
predict.Random.Forest
```

```
##      [1] 0 1 1 0 1 0 1 0 0 0 0 1 0 0 0 0 0 0 0 1 0 1 0 0 1 1 0 1 1 0 0 1 1 0 1
##     [38] 1 1 1 1 1 0 1 1 1 1 1 0 0 1 1 1 1 1 1 0 0 0 0 0 1 0 0 0 0 0 0 1 1 0 1 0
##     [75] 1 1 0 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 1 1 0
##    [112] 0 1 0 1 0 1 0 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 1 0 0 1 0 0 0 1 0
##    [149] 0 0 1 1 0
## Levels: 0 1
```

Confusion Matrix

```
confusionMatrix(predict.Random.Forest,diabetes.test$Outcome)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction  0   1
##           0 79 14
##           1 17 43
##
##              Accuracy : 0.7974
##              95% CI : (0.7249, 0.858)
##      No Information Rate : 0.6275
##      P-Value [Acc > NIR] : 4.322e-06
##
##              Kappa : 0.5712
##
##  Mcnemar's Test P-Value : 0.7194
##
##              Sensitivity : 0.8229
##              Specificity : 0.7544
##              Pos Pred Value : 0.8495
##              Neg Pred Value : 0.7167
##              Prevalence : 0.6275
##              Detection Rate : 0.5163
##      Detection Prevalence : 0.6078
##              Balanced Accuracy : 0.7887
##
##              'Positive' Class : 0
##
```

Accuracy is 79% the percentage of correctly classified instances out of all instances.

Sensitivity is 82% which is the true positive rate. It is the number instances from the positive class that actually predicted diabetes outcome correctly.

Specificity is 75% which is the true negative rate. Is the number of instances from the negative class that actually predicted diabetes outcome correctly.

##Model2: LOGISTIC REGRESSION

Build a Logistic Regression Model based on variables

```
model_glm<-glm(Outcome~Pregnancies+Glucose+BMI+SkinThickness+Insulin+DiabetesPedigreeFunction+Age,data=
summary(model_glm)
```

```
##
## Call:
## glm(formula = Outcome ~ Pregnancies + Glucose + BMI + SkinThickness +
##      Insulin + DiabetesPedigreeFunction + Age, family = binomial,
##      data = diabetes.train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4915  -0.7566  -0.4448   0.7685   2.7520
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -8.522945    0.781303  -10.909  < 2e-16 ***
## Pregnancies     0.117029    0.035605   3.287  0.00101 **
## Glucose         0.034151    0.004088   8.354  < 2e-16 ***
## BMI             0.083839    0.016970   4.940  7.8e-07 ***
## SkinThickness   0.002600    0.007609   0.342  0.73257
## Insulin        -0.001228    0.001035  -1.186  0.23556
## DiabetesPedigreeFunction 0.696008    0.329683   2.111  0.03476 *
## Age             0.004259    0.010097   0.422  0.67316
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 790.97  on 614  degrees of freedom
## Residual deviance: 600.13  on 607  degrees of freedom
## AIC: 616.13
```

```
##
## Number of Fisher Scoring iterations: 5
```

```
#Test model on the test data
```

```
predict_test <- predict(model_glm,newdata = diabetes.test,type = "response")
```

```
predict_test
```

```
##          2          5          9         22         27         29
## 0.056185761 0.800102880 0.686114407 0.319754341 0.727438069 0.534731612
##          31          34          43          51          53          65
## 0.378948192 0.059500896 0.146060817 0.052225398 0.076230366 0.331323201
##          66          75          80          81          84          87
## 0.144577490 0.070577603 0.119733829 0.098864295 0.064173446 0.535652118
##          88          107         110         112         118         132
## 0.208838937 0.047861557 0.136093991 0.643587631 0.131585148 0.599750467
##          135          143          144          148          150          153
## 0.062577753 0.172136595 0.358975069 0.259668801 0.061256722 0.835132750
##          160          162          170          176          179          184
## 0.960743187 0.311564100 0.168453624 0.866591704 0.740948802 0.052331901
##          187          188          193          194          205          212
## 0.783749799 0.436246095 0.667034277 0.911503724 0.303052206 0.643164456
##          217          236          237          247          248          251
## 0.309709860 0.868118359 0.873395360 0.446130537 0.758213241 0.312043458
##          252          257          260          261          267          271
## 0.227354619 0.224307582 0.869566689 0.763842545 0.497260836 0.724785144
##          284          296          304          308          309          316
## 0.658861254 0.699819542 0.665815363 0.151689157 0.333696540 0.208610586
##          325          335          336          343          348          349
## 0.232309270 0.051779490 0.771635912 0.005129434 0.117942215 0.062986623
##          350          351          354          356          358          359
## 0.017657260 0.252277323 0.071251393 0.756346902 0.803838573 0.343307692
##          364          374          375          376          381          388
## 0.691307764 0.178667448 0.381662932 0.781826461 0.179670982 0.520222256
##          391          392          395          403          409          410
## 0.120755403 0.871498659 0.687872228 0.496937265 0.909490394 0.740249555
##          411          421          426          427          432          434
## 0.341955075 0.454487175 0.803461556 0.006508378 0.108985579 0.239407292
##          450          466          470          472          473          479
## 0.169493548 0.104178301 0.852023158 0.320813430 0.296941641 0.304339796
##          480          487          498          502          504          509
## 0.363800220 0.433679512 0.071343356 0.142325562 0.260100275 0.105095465
##          512          513          514          516          522          527
## 0.129742119 0.124897208 0.080670176 0.566526217 0.288128461 0.034852706
##          546          549          551          559          569          584
## 0.872644008 0.604571584 0.136693154 0.628508903 0.535462636 0.350699886
##          585          590          591          592          597          599
## 0.284885147 0.019523159 0.805059356 0.273170089 0.108514586 0.692216785
##          607          608          610          612          618          620
## 0.863857982 0.040210866 0.069224099 0.713769699 0.018176839 0.176203044
##          624          638          643          644          654          662
## 0.196896858 0.128535652 0.502364324 0.111333765 0.180605016 0.959397681
```

```
##           665           667           668           672           673           674
## 0.348465240 0.532522141 0.285433365 0.083435678 0.160307600 0.791702281
##           681           682           683           688           692           695
## 0.017894359 0.835559098 0.228205843 0.111029553 0.909445534 0.047168847
##           705           708           709           720           723           740
## 0.142966474 0.230032075 0.764155031 0.252417596 0.372921539 0.226513485
##           749           756           765
## 0.819805592 0.471312112 0.350828147
```

```
predict_test <- ifelse(predict_test > 0.5,1,0)
```

Confusion Matrix


```
confusionMatrix(factor(diabetes.test$Outcome), factor(predict_test))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 83 13
##           1 17 40
##
##           Accuracy : 0.8039
##           95% CI : (0.7321, 0.8636)
##           No Information Rate : 0.6536
##           P-Value [Acc > NIR] : 3.3e-05
##
##           Kappa : 0.5745
##
## Mcnemar's Test P-Value : 0.5839
##
##           Sensitivity : 0.8300
##           Specificity : 0.7547
##           Pos Pred Value : 0.8646
##           Neg Pred Value : 0.7018
##           Prevalence : 0.6536
##           Detection Rate : 0.5425
##           Detection Prevalence : 0.6275
##           Balanced Accuracy : 0.7924
##
##           'Positive' Class : 0
##
```

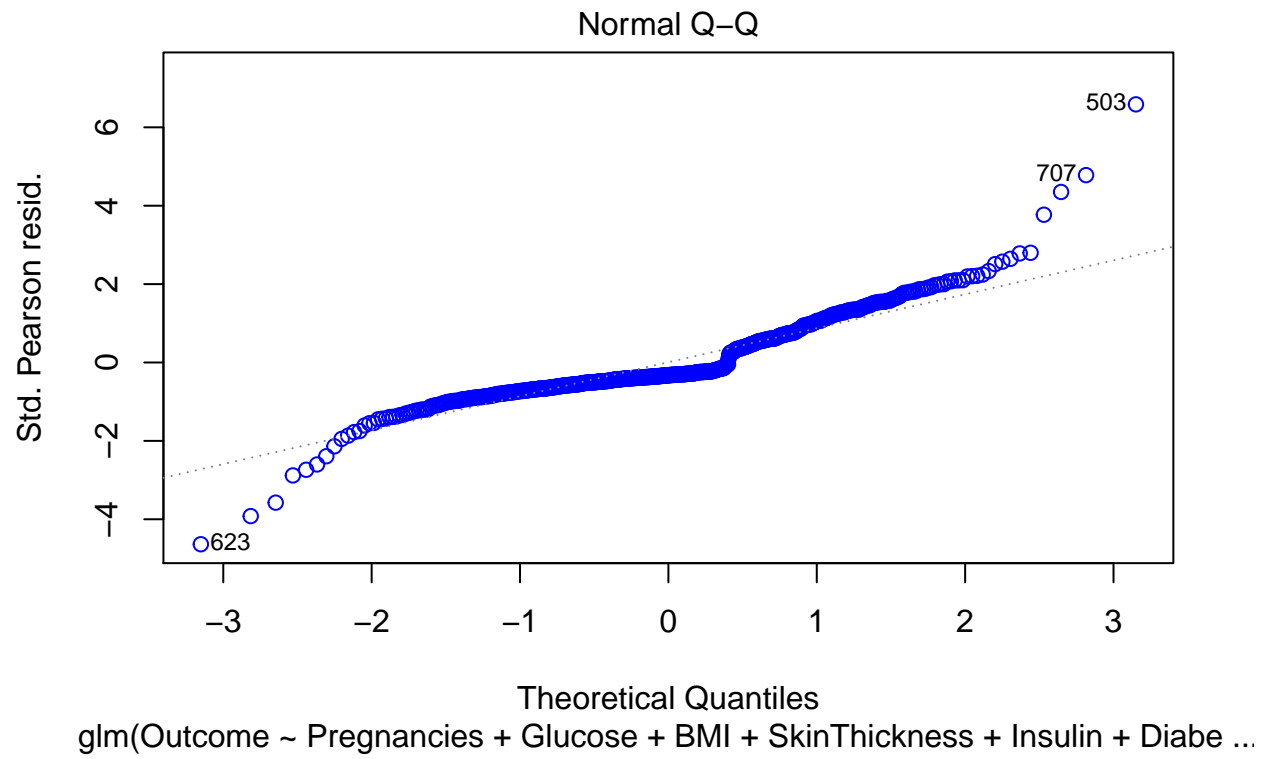
Applying the Logistic Regression Algorithm, the Accuracy is 80%, sensitivity is 83% and specificity is 75%

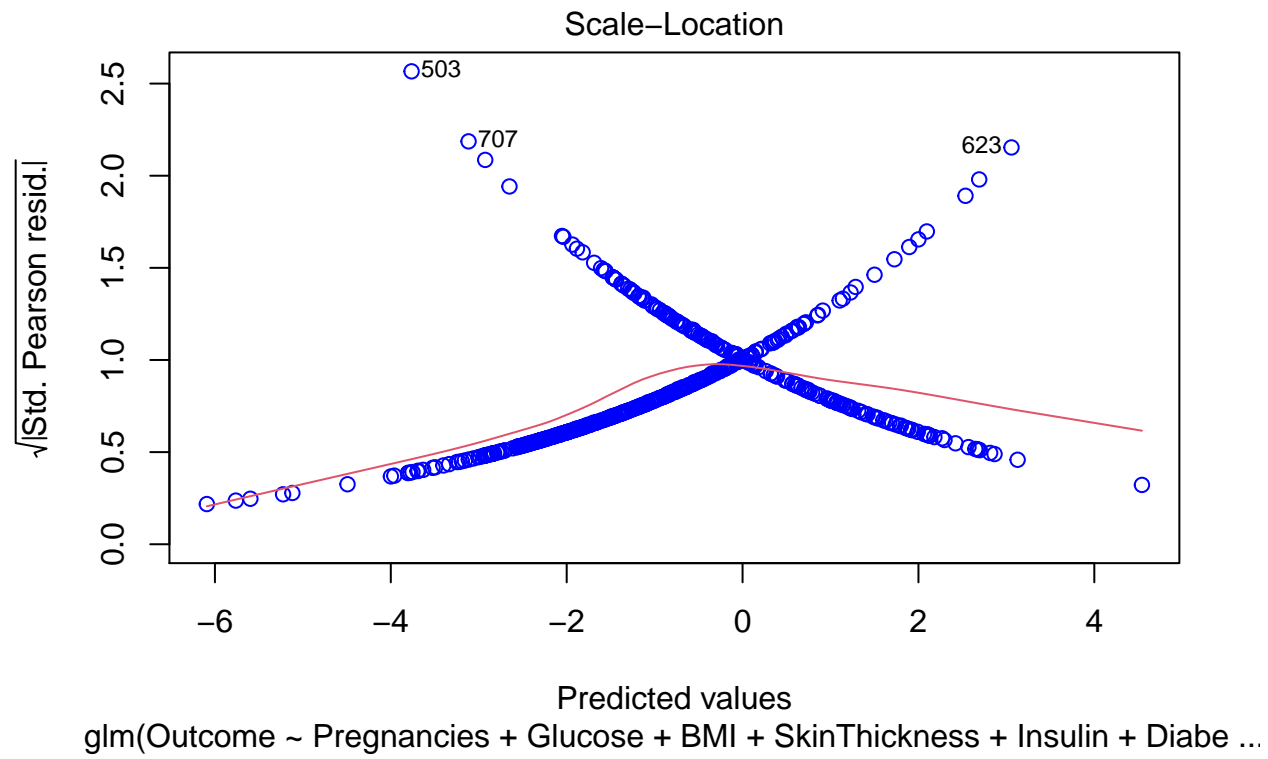
Plotting ROC curve of Logistic Regression Model

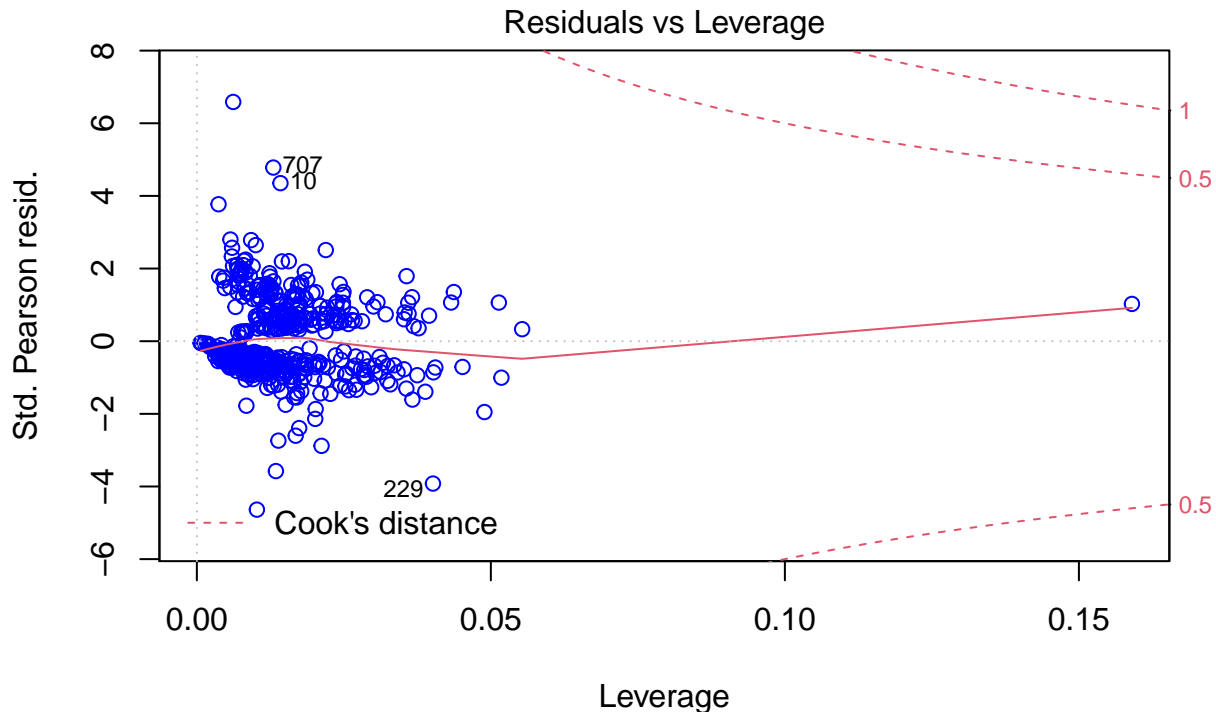
```
plot(model_glm, col='blue')
```



glm(Outcome ~ Pregnancies + Glucose + BMI + SkinThickness + Insulin + Diabe ...







glm(Outcome ~ Pregnancies + Glucose + BMI + SkinThickness + Insulin + Diabe ...

The Z score is a test of statistical significance that helps you decide whether or not to reject the null hypothesis. It also gives you an idea of how far from the mean a data point is. Looking at the summary, we can see which variables are significant by comparing the p-values. P-values with '***' next to them are significant and play a role in whether a subject has diabetes or not.

Plot the results of the two machine learning models


```
# prepare training scheme
control <- trainControl(method="repeatedcv", number=10, repeats=3)
set.seed(7)
fit.rf <- train(Outcome~., data=diabetes, method="rf", trControl=control)
set.seed(7)
fit.glm <- train(Outcome~., data=diabetes, method="glm", trControl=control)
# collect resamples
results <- resamples(list(LR=fit.glm, RF=fit.rf))
```

Summarize the distributions



```
summary(results)
```

```
##  
## Call:  
## summary.resamples(object = results)  
##  
## Models: LR, RF  
## Number of resamples: 30  
##  
## Accuracy  
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's  
## LR 0.7272727 0.7508117 0.7662338 0.7812657 0.8000256 0.9078947    0  
## RF 0.6842105 0.7305195 0.7597403 0.7638528 0.8019481 0.8421053    0  
##  
## Kappa  
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's  
## LR 0.3513839 0.4168485 0.4662541 0.4931161 0.5391907 0.7812500    0  
## RF 0.2951613 0.3778304 0.4640696 0.4630809 0.5447483 0.6426332    0
```

```
# boxplots of results  
bwplot(results)
```




```
# dot plots of results
dotplot(results)
```




The distributions are summarized in terms of the percentiles. The distributions are summarized as box plots and finally the distributions are summarized as dot plots.

```
# density plots of accuracy
scales <- list(x=list(relation="free"), y=list(relation="free"))
densityplot(results, scales=scales, pch = "|")
```



xyplot plots to compare models

```
xyplot(results, models=c("LR", "RF"))
```



Difference in model predictions

```
diffs <- diff(results)

# summarize p-values for pair-wise comparisons
summary(diffs)
```

```
##
## Call:
## summary.diff.resamples(object = diffs)
##
## p-value adjustment: bonferroni
## Upper diagonal: estimates of the difference
## Lower diagonal: p-value for H0: difference = 0
##
## Accuracy
##      LR      RF
## LR      0.01741
## RF 0.01599
##
## Kappa
##      LR      RF
## LR      0.03004
```

RF 0.07278

Results:

**Context:* The goal of the project was to predict diabetes among PIMA women using variables such as BMI, Glucose, Pregnancies, Pedigree, Age, Family, Insulin, Skin Thickness.

Problem: The problem was to predict diabetes and the Machine Learning Models in this project were able to effectively predict the outcome.

Solution: We can interpret from the p-values of our Models that BMI and glucose are the biggest factors in determining whether members of the PIMA Indian have diabetes. But these two alone are not sufficient enough to accurately predict the outcome.

Findings: Accuracy is 79% the percentage of correctly classified instances out of all instances using Random Forest Model and applying the Logistic Regression Algorithm, the Accuracy is 80%.

Limitations: The sample size being not large enough to validate accurate predictions could be considered a limitation. A data set of at least 100,000 or even a million observations would be ideal for accurate predictions. The sample data comprising of only women, is also a limitation to universally apply the prediction algorithms.

Conclusions:

Based on the concepts learned in the data science course series, all aspects of building an effective model to predict the outcome has been implemented. However, the larger the data set more accurate the predicted outcome would be. Patterns were established using data exploration and validation. The project shows us what are the most important factors that influence a person to have diabetes. Predictive models improve prediction performance but they don't provide outstanding results. Maybe other Machine learning models can be tried to see how it influences the results. The patterns identified using Data exploration methods were validated using the modeling techniques employed. Based on the above modelling, Logistic Regression is the best predictive model to determine if there is a possibility of diabetic outcome in a person.

#Reference:

1. University of Chicago Press Journals, <https://www.journals.uchicago.edu/doi/full/10.1086/693853?mobileUi=0&>
2. <https://fderyckel.github.io/machinelearningwithr/logistic.html>
3. <http://www.joyofdata.de/blog/illustrated-guide-to-roc-and-auc/>
4. <https://github.com/joyofdata/joyofdata-articles/tree/master/roc-auc>
5. <https://machinelearningmastery.com/compare-the-performance-of-machine-learning-algorithms-in-r/>
6. <https://www.kaggle.com/nileshvarshney1/pima-indians-diabetes>